

```

# First, make sure to import the yfinance library
import yfinance as yf

# Create a Ticker object for Nifty 50 (the correct symbol might be
#"^NSEI")
nifty50 = yf.Ticker("^NSEI") # Use the correct ticker symbol for
Nifty 50

# Now you can use the history method
nifty50_data = nifty50.history(period="max")

nifty50_data

```

		Open	High	Low	\
Date					
2007-09-17 00:00:00+05:30		4518.450195	4549.049805	4482.850098	
2007-09-18 00:00:00+05:30		4494.100098	4551.799805	4481.549805	
2007-09-19 00:00:00+05:30		4550.250000	4739.000000	4550.250000	
2007-09-20 00:00:00+05:30		4734.850098	4760.850098	4721.149902	
2007-09-21 00:00:00+05:30		4752.950195	4855.700195	4733.700195	
...		
2025-12-08 00:00:00+05:30		26159.800781	26178.699219	25892.250000	
2025-12-09 00:00:00+05:30		25867.099609	25923.650391	25728.000000	
2025-12-10 00:00:00+05:30		25864.050781	25947.650391	25734.550781	
2025-12-11 00:00:00+05:30		25771.400391	25922.800781	25693.250000	
2025-12-12 00:00:00+05:30		25971.199219	26057.599609	25938.449219	
		Close	Volume	Dividends	Stock
Splits					
Date					
2007-09-17 00:00:00+05:30		4494.649902	0	0.0	
0.0					
2007-09-18 00:00:00+05:30		4546.200195	0	0.0	
0.0					
2007-09-19 00:00:00+05:30		4732.350098	0	0.0	
0.0					
2007-09-20 00:00:00+05:30		4747.549805	0	0.0	
0.0					
2007-09-21 00:00:00+05:30		4837.549805	0	0.0	
0.0					
...		
...		
2025-12-08 00:00:00+05:30		25960.550781	320700	0.0	
0.0					
2025-12-09 00:00:00+05:30		25839.650391	276000	0.0	
0.0					
2025-12-10 00:00:00+05:30		25758.000000	207900	0.0	
0.0					
2025-12-11 00:00:00+05:30		25898.550781	206100	0.0	

```
0.0
2025-12-12 00:00:00+05:30  26046.949219          0        0.0
0.0

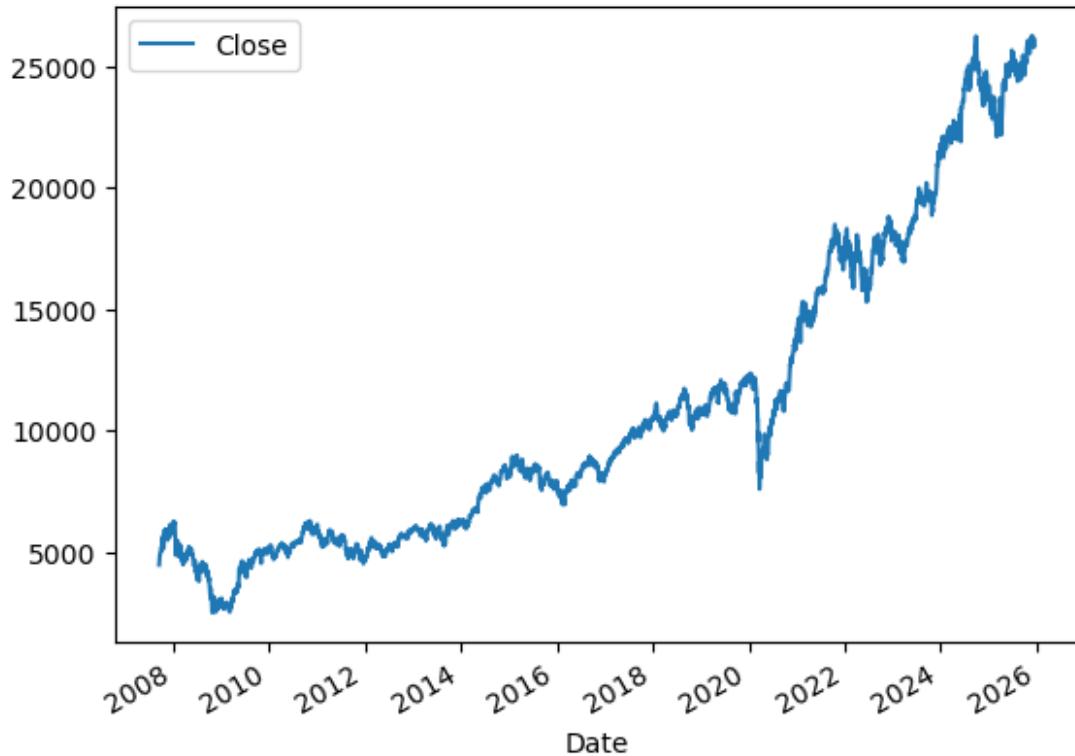
[4475 rows x 7 columns]

nifty50_data.index

DatetimeIndex(['2007-09-17 00:00:00+05:30', '2007-09-18
00:00:00+05:30',
               '2007-09-19 00:00:00+05:30', '2007-09-20
00:00:00+05:30',
               '2007-09-21 00:00:00+05:30', '2007-09-24
00:00:00+05:30',
               '2007-09-25 00:00:00+05:30', '2007-09-26
00:00:00+05:30',
               '2007-09-27 00:00:00+05:30', '2007-09-28
00:00:00+05:30',
               ...
               '2025-12-01 00:00:00+05:30', '2025-12-02
00:00:00+05:30',
               '2025-12-03 00:00:00+05:30', '2025-12-04
00:00:00+05:30',
               '2025-12-05 00:00:00+05:30', '2025-12-08
00:00:00+05:30',
               '2025-12-09 00:00:00+05:30', '2025-12-10
00:00:00+05:30',
               '2025-12-11 00:00:00+05:30', '2025-12-12
00:00:00+05:30'],
              dtype='datetime64[ns, Asia/Kolkata]', name='Date',
              length=4475, freq=None)

# cleaning and visualizing our stock market
nifty50_data.plot.line(y="Close",use_index=True)

<Axes: xlabel='Date'>
```



```
# Check if columns exist before deleting them
if "Stock Splits" in nifty50_data.columns:
    del nifty50_data["Stock Splits"]

if "Dividends" in nifty50_data.columns:
    del nifty50_data["Dividends"]

# Alternative approach using drop method with errors='ignore'
# nifty50_data = nifty50_data.drop(["Stock Splits", "Dividends"], axis=1, errors='ignore')

nifty50_data
```

Date	Open	High	Low	\
2007-09-17 00:00:00+05:30	4518.450195	4549.049805	4482.850098	
2007-09-18 00:00:00+05:30	4494.100098	4551.799805	4481.549805	
2007-09-19 00:00:00+05:30	4550.250000	4739.000000	4550.250000	
2007-09-20 00:00:00+05:30	4734.850098	4760.850098	4721.149902	
2007-09-21 00:00:00+05:30	4752.950195	4855.700195	4733.700195	
...
2025-12-08 00:00:00+05:30	26159.800781	26178.699219	25892.250000	
2025-12-09 00:00:00+05:30	25867.099609	25923.650391	25728.000000	
2025-12-10 00:00:00+05:30	25864.050781	25947.650391	25734.550781	
2025-12-11 00:00:00+05:30	25771.400391	25922.800781	25693.250000	
2025-12-12 00:00:00+05:30	25971.199219	26057.599609	25938.449219	

		Close	Volume
Date			
2007-09-17 00:00:00+05:30		4494.649902	0
2007-09-18 00:00:00+05:30		4546.200195	0
2007-09-19 00:00:00+05:30		4732.350098	0
2007-09-20 00:00:00+05:30		4747.549805	0
2007-09-21 00:00:00+05:30		4837.549805	0
...	
2025-12-08 00:00:00+05:30		25960.550781	320700
2025-12-09 00:00:00+05:30		25839.650391	276000
2025-12-10 00:00:00+05:30		25758.000000	207900
2025-12-11 00:00:00+05:30		25898.550781	206100
2025-12-12 00:00:00+05:30		26046.949219	0

[4475 rows x 5 columns]

```
import numpy as np
import pandas as pd

df = nifty50_data.copy()

# Daily returns
df["Return"] = df["Close"].pct_change()

# Rolling volatility
df["Volatility"] = df["Return"].rolling(5).std()

# Moving averages
df["MA10"] = df["Close"].rolling(10).mean()
df["MA20"] = df["Close"].rolling(20).mean()

df.dropna(inplace=True)
df.head()
```

		Open	High	Low
Close \ Date				
2007-10-15 00:00:00+05:30	5428.350098	5682.649902	5419.899902	
5670.399902				
2007-10-16 00:00:00+05:30	5670.649902	5708.350098	5578.450195	
5668.049805				
2007-10-17 00:00:00+05:30	5658.899902	5658.899902	5107.299805	
5559.299805				
2007-10-18 00:00:00+05:30	5551.100098	5736.799805	5269.649902	
5351.000000				
2007-10-19 00:00:00+05:30	5360.350098	5390.850098	5101.750000	
5215.299805				

Volume Return Volatility

```

MA10 \
Date

2007-10-15 00:00:00+05:30      0  0.044609    0.026320  5315.155029
2007-10-16 00:00:00+05:30      0 -0.000414    0.023370  5375.064990
2007-10-17 00:00:00+05:30      0 -0.019186    0.026431  5409.914990
2007-10-18 00:00:00+05:30      0 -0.037469    0.031179  5424.150000
2007-10-19 00:00:00+05:30      0 -0.025360    0.032091  5427.094971

MA20
Date

2007-10-15 00:00:00+05:30  5067.165015
2007-10-16 00:00:00+05:30  5125.835010
2007-10-17 00:00:00+05:30  5176.489990
2007-10-18 00:00:00+05:30  5207.422485
2007-10-19 00:00:00+05:30  5230.809985

# Median volatility as threshold
vol_threshold = df["Volatility"].median()

df["Target"] = (df["Volatility"].shift(-1) >
vol_threshold).astype(int)

df.dropna(inplace=True)

df[["Volatility", "Target"]]

          Volatility  Target
Date
2007-10-15 00:00:00+05:30    0.026320     1
2007-10-16 00:00:00+05:30    0.023370     1
2007-10-17 00:00:00+05:30    0.026431     1
2007-10-18 00:00:00+05:30    0.031179     1
2007-10-19 00:00:00+05:30    0.032091     1
...
2025-12-08 00:00:00+05:30    0.005745     0
2025-12-09 00:00:00+05:30    0.005617     0
2025-12-10 00:00:00+05:30    0.005670     0
2025-12-11 00:00:00+05:30    0.006423     0
2025-12-12 00:00:00+05:30    0.006387     0

[4456 rows x 2 columns]

features = ["Return", "MA10", "MA20", "Volatility"]

```

```

X = df[features]
y = df["Target"]

from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, shuffle=False
)

from sklearn.ensemble import RandomForestClassifier

model = RandomForestClassifier(
    n_estimators=200,
    max_depth=5,
    random_state=42
)

model.fit(X_train, y_train)

RandomForestClassifier(max_depth=5, n_estimators=200, random_state=42)

from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix

y_pred = model.predict(X_test)

print("Accuracy:", accuracy_score(y_test, y_pred))
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test,
y_pred))

Accuracy: 0.8778026905829597

Confusion Matrix:
[[599  62]
 [ 47 184]]

Classification Report:
              precision    recall  f1-score   support

             0       0.93     0.91     0.92      661
             1       0.75     0.80     0.77      231

        accuracy                           0.88      892
       macro avg       0.84     0.85     0.84      892
    weighted avg       0.88     0.88     0.88      892

latest_data = X.iloc[-1:].values
prediction = model.predict(latest_data)[0]

if prediction == 1:

```

```

    print("⚠️ High volatility expected tomorrow")
else:
    print("⚠️ Low volatility expected tomorrow")

⚠️ Low volatility expected tomorrow

C:\Users\chait\AppData\Local\Programs\Orange\Lib\site-packages\
sklearn\base.py:465: UserWarning: X does not have valid feature names,
but RandomForestClassifier was fitted with feature names
    warnings.warn(
df = nifty50_data.copy()

df["Target"] = df["Close"].shift(-1)

df.dropna(inplace=True)

df[["Close", "Target"]]

[4474 rows x 2 columns]

# Returns
df["Return"] = df["Close"].pct_change()

# Moving averages
df["MA5"] = df["Close"].rolling(5).mean()
df["MA10"] = df["Close"].rolling(10).mean()

# Volatility
df["Volatility"] = df["Return"].rolling(5).std()

df.dropna(inplace=True)

features = ["Close", "Return", "MA5", "MA10", "Volatility"]

X = df[features]
y = df["Target"]

```

```

from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, shuffle=False
)

from sklearn.ensemble import RandomForestRegressor

model = RandomForestRegressor(
    n_estimators=300,
    max_depth=6,
    random_state=42
)

model.fit(X_train, y_train)

RandomForestRegressor(max_depth=6, n_estimators=300, random_state=42)

from sklearn.metrics import mean_absolute_error, mean_squared_error

y_pred = model.predict(X_test)

mae = mean_absolute_error(y_test, y_pred)
rmse = mean_squared_error(y_test, y_pred, squared=False)

print("MAE:", mae)
print("RMSE:", rmse)

MAE: 3405.4683136570684
RMSE: 4405.556024129921

latest_features = X.iloc[-1:].values
tomorrow_close = model.predict(latest_features)[0]

print(" Predicted NIFTY 50 Close Tomorrow:", round(tomorrow_close, 2))

 Predicted NIFTY 50 Close Tomorrow: 18196.16

C:\Users\chait\AppData\Local\Programs\Orange\Lib\site-packages\
sklearn\base.py:465: UserWarning: X does not have valid feature names,
but RandomForestRegressor was fitted with feature names
warnings.warn(
df = nifty50_data.copy()

df["Return"] = df["Close"].pct_change()
df["Target"] = df["Return"].shift(-1)

df.dropna(inplace=True)

```

```

df["MA5"] = df["Close"].rolling(5).mean()
df["MA10"] = df["Close"].rolling(10).mean()
df["Volatility"] = df["Return"].rolling(5).std()

df.dropna(inplace=True)

features = ["Return", "MA5", "MA10", "Volatility"]

X = df[features]
y = df["Target"]

from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y, test_size=0.2, shuffle=False
)

from sklearn.ensemble import RandomForestRegressor

model = RandomForestRegressor(
    n_estimators=300,
    max_depth=6,
    random_state=42
)

model.fit(X_train, y_train)

RandomForestRegressor(max_depth=6, n_estimators=300, random_state=42)

from sklearn.metrics import mean_absolute_error, mean_squared_error

y_pred = model.predict(X_test)

print("MAE (Return):", mean_absolute_error(y_test, y_pred))
print("RMSE (Return):", mean_squared_error(y_test, y_pred,
squared=False))

MAE (Return): 0.005811790142321782
RMSE (Return): 0.007928511346090988

latest_return_pred = model.predict(X_scaled[-1:].reshape(1, -1))[0]

last_close = df["Close"].iloc[-1]
predicted_close = last_close * (1 + latest_return_pred)

print(" Predicted NIFTY Close Tomorrow:", round(predicted_close, 2))

```

```

□ Predicted NIFTY Close Tomorrow: 25900.51

#On average, your model's next-day return prediction error is less
than 1%.
```

```

# Predict next-day return
predicted_return = model.predict(X_scaled[-1:]).reshape(1, -1))[0]

# Convert return to price
last_close = df["Close"].iloc[-1]
predicted_close = last_close * (1 + predicted_return)

print("Last Close:", round(last_close, 2))
print("Predicted Return:", round(predicted_return * 100, 3), "%")
print("□ Predicted NIFTY Close Tomorrow:", round(predicted_close, 2))

Last Close: 25898.55
Predicted Return: 0.019 %
□ Predicted NIFTY Close Tomorrow: 25903.51
```

```

from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, shuffle=False
)

from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()

X_train_scaled = scaler.fit_transform(X_train) # FIT ONLY ON TRAIN
X_test_scaled = scaler.transform(X_test) # TRANSFORM TEST

model.fit(X_train_scaled, y_train)

RandomForestRegressor(max_depth=6, n_estimators=300, random_state=42)

y_pred = model.predict(X_test_scaled)

print("MAE:", mean_absolute_error(y_test, y_pred))
print("RMSE:", mean_squared_error(y_test, y_pred, squared=False))

MAE: 0.005811790142321782
RMSE: 0.007928511346090988

# Scale test data correctly
X_test_scaled = scaler.transform(X_test)

# Predict returns
y_test_pred = model.predict(X_test_scaled)
```

```

results = pd.DataFrame({
    "Actual_Return": y_test.values,
    "Predicted_Return": y_test_pred
}, index=y_test.index)

results

      Actual_Return  Predicted_Return
Date
2022-05-05 00:00:00+05:30     -0.016268      0.000061
2022-05-06 00:00:00+05:30     -0.006666      0.000049
2022-05-09 00:00:00+05:30     -0.003791      0.000049
2022-05-10 00:00:00+05:30     -0.004492      0.000049
2022-05-11 00:00:00+05:30     -0.022212      0.000049
...
2025-12-05 00:00:00+05:30     -0.008627      0.000094
2025-12-08 00:00:00+05:30     -0.004657     -0.000049
2025-12-09 00:00:00+05:30     -0.003160     -0.000049
2025-12-10 00:00:00+05:30      0.005457     -0.000049
2025-12-11 00:00:00+05:30      0.005730      0.000076

[893 rows x 2 columns]

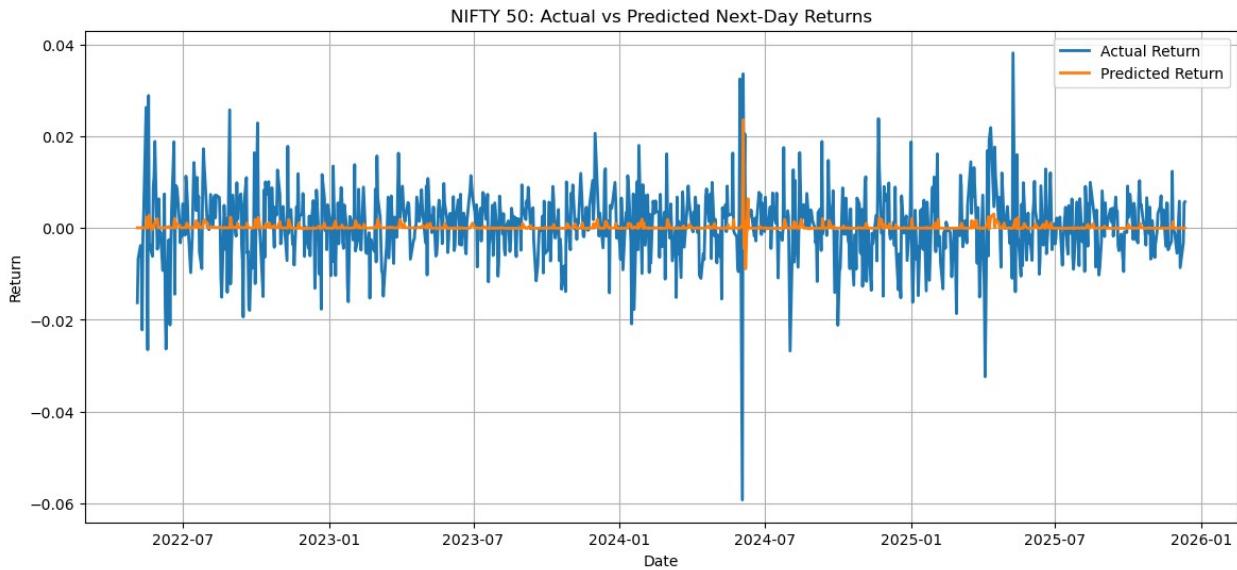
import matplotlib.pyplot as plt

plt.figure(figsize=(14,6))
plt.plot(results.index, results["Actual_Return"], label="Actual Return", linewidth=2)
plt.plot(results.index, results["Predicted_Return"], label="Predicted Return", linewidth=2)

plt.title("NIFTY 50: Actual vs Predicted Next-Day Returns")
plt.xlabel("Date")
plt.ylabel("Return")
plt.legend()
plt.grid(True)

plt.show()

```



```

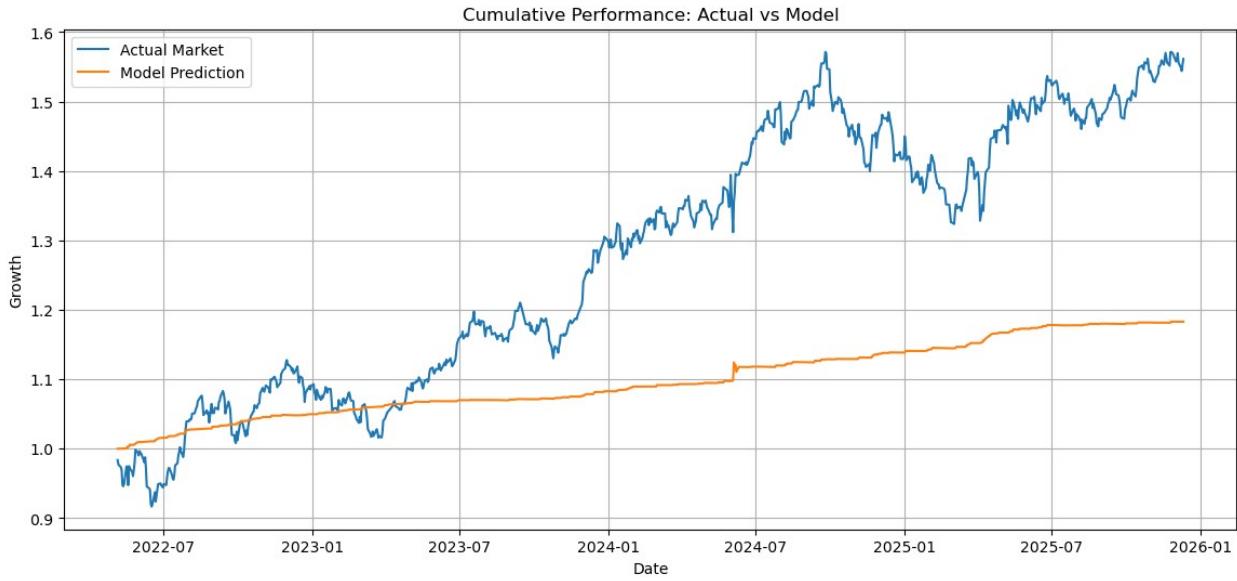
results["Actual_Cumulative"] = (1 +
results["Actual_Return"]).cumprod()
results["Predicted_Cumulative"] = (1 +
results["Predicted_Return"]).cumprod()

plt.figure(figsize=(14,6))
plt.plot(results.index, results["Actual_Cumulative"], label="Actual Market")
plt.plot(results.index, results["Predicted_Cumulative"], label="Model Prediction")

plt.title("Cumulative Performance: Actual vs Model")
plt.xlabel("Date")
plt.ylabel("Growth")
plt.legend()
plt.grid(True)

plt.show()

```



```
# Direction: 1 = Up, 0 = Down
results["Actual_Direction"] = (results["Actual_Return"] > 0).astype(int)
results["Predicted_Direction"] = (results["Predicted_Return"] > 0).astype(int)

direction_accuracy = (
    results["Actual_Direction"] == results["Predicted_Direction"]
).mean()

print("Directional Accuracy:", round(direction_accuracy * 100, 2),
"%")

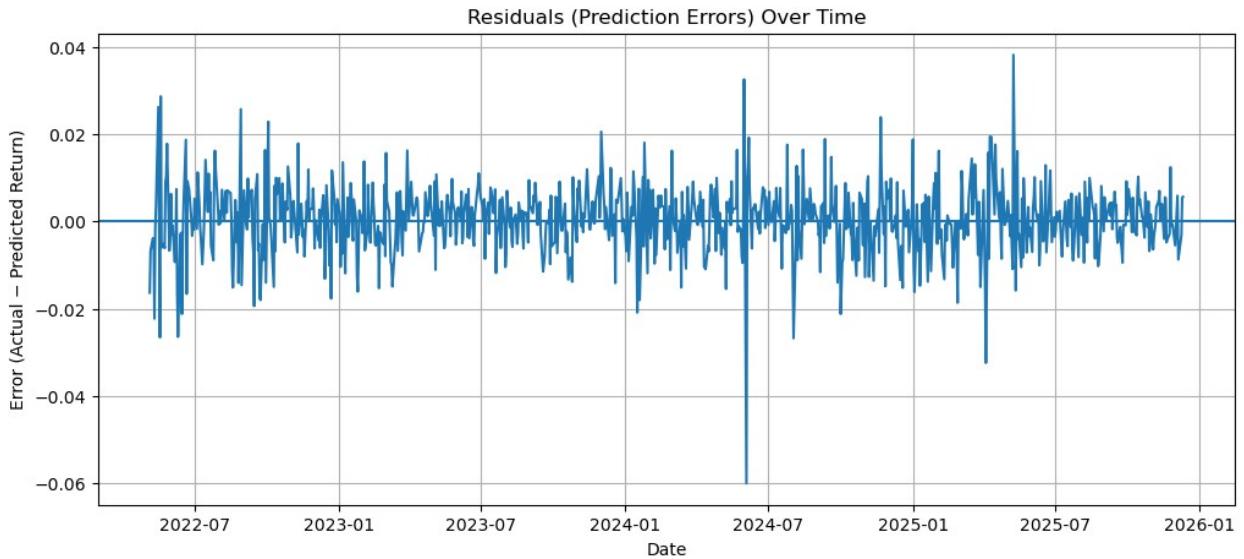
Directional Accuracy: 50.06 %

# random

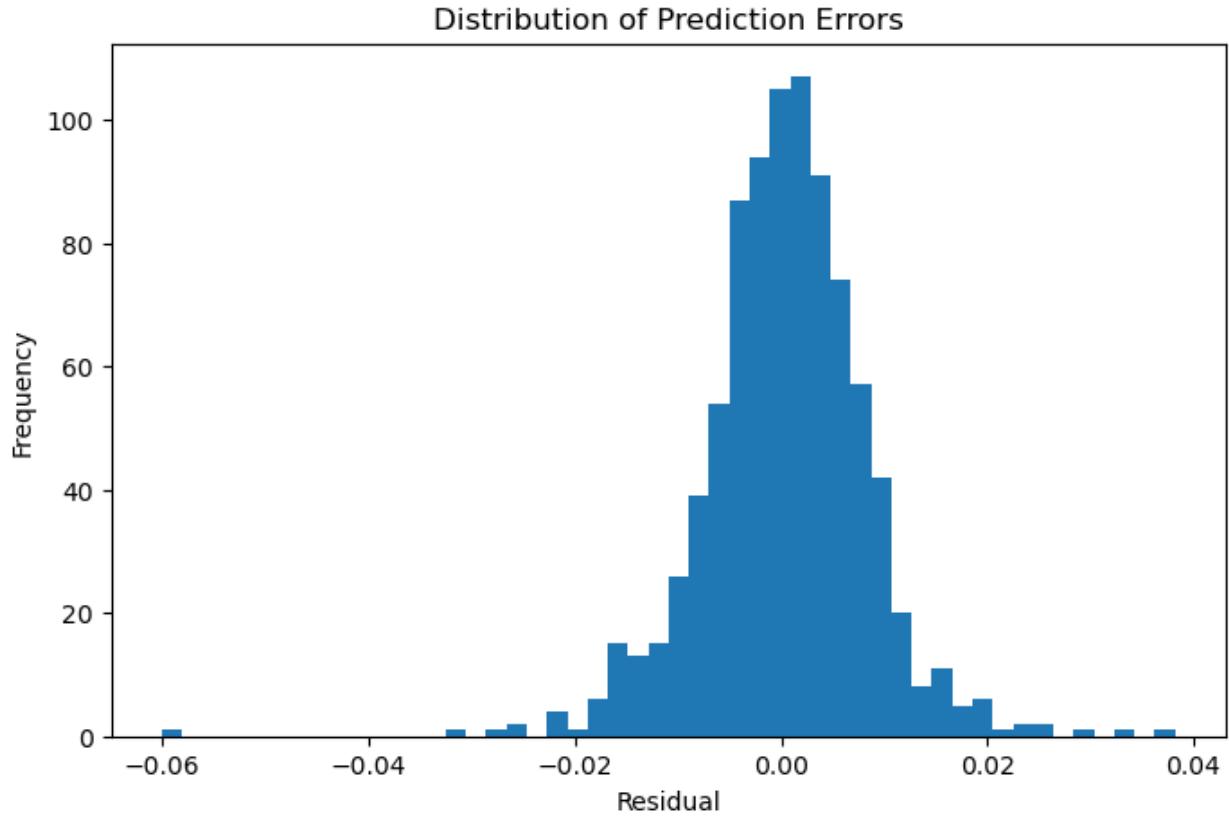
results["Residual"] = results["Actual_Return"] -
results["Predicted_Return"]

plt.figure(figsize=(12,5))
plt.plot(results.index, results["Residual"])
plt.axhline(0)

plt.title("Residuals (Prediction Errors) Over Time")
plt.xlabel("Date")
plt.ylabel("Error (Actual - Predicted Return)")
plt.grid(True)
plt.show()
```



```
plt.figure(figsize=(8,5))
plt.hist(results["Residual"], bins=50)
plt.title("Distribution of Prediction Errors")
plt.xlabel("Residual")
plt.ylabel("Frequency")
plt.show()
```



```
import pandas as pd

importance = pd.Series(
    model.feature_importances_,
    index=X.columns
).sort_values(ascending=False)

importance

Volatility      0.360744
Return          0.268574
MA10            0.199126
MA5             0.171556
dtype: float64

importance.plot(kind="bar", figsize=(8,4), title="Feature Importance")
plt.show()
```

